

MODEL AND VARIABLE SELECTION

ARTHUR ALLIGNOL

INTRODUCTION

WHY DO WE WANT TO CHOOSE A MODEL?

Explanation:

Occam's razor applies and we want

an explanation that is as simple as possible, but not simpler

Prediction:

For prediction, we want a model that predicts well (!) All that matters is that it works, the model doesn't necessarily need to be interpretable

BIAS-VARIANCE TRADE-OFF

Let $Y = f(X) + \varepsilon$, with $E(\varepsilon) = 0$ and $\text{var}\varepsilon = \sigma_\varepsilon^2$. The error of a regression fit $\hat{f}(X)$ at an input point x_0 is (using squared error)

$$\begin{aligned}\text{Err}(x_0) &= E[(Y - \hat{f}(x_0))^2 | X = x_0] \\ &= \sigma_\varepsilon^2 + [E\hat{f}(x_0) - f(x_0)]^2 + E[\hat{f}(x_0) - E\hat{f}(x_0)]^2 \\ &= \sigma_\varepsilon^2 + \text{Bias}^2(f(\hat{x}_0)) + \text{var}(\hat{f}(x_0)) \\ &= \text{Irreducible error} + \text{Bias}^2 + \text{Variance}\end{aligned}$$

Typically, the more complex we make the model \hat{f} , the lower the (squared) error but the higher the variance

MODEL SELECTION IS DIFFICULT

```
set.seed(442)
n <- 200
p <- 100
y <- rnorm(n)
x <- matrix(rnorm(n * p), nrow = n)
df <- data.frame(y, x)
mdl <- lm(y ~ ., df)
```

MODEL SELECTION IS DIFFICULT

```
stars <- 1 + which(coefficients(summary mdl))[-1, 4] < 0.05)
mdl.2 <- lm(y ~ ., df[, c(1, stars)])
summary(mdl.2)
```

Call:

```
lm(formula = y ~ ., data = df[, c(1, stars)])
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.56045	-0.73379	0.06531	0.63945	2.83769

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	0.00228	0.07710	0.030	0.9764	
X4	-0.15414	0.07373	-2.090	0.0379	*
X14	0.09399	0.07275	1.292	0.1979	
X40	-0.09558	0.07266	-1.315	0.1899	
X63	-0.17910	0.08089	-2.214	0.0280	*
X65	-0.19684	0.07736	-2.544	0.0117	*
X84	-0.12703	0.07221	-1.759	0.0801	.
X99	0.05094	0.07675	0.664	0.5077	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.048 on 192 degrees of freedom

Multiple R-squared: 0.09739 Adjusted R-squared: 0.06119

MODEL SELECTION IS DIFFICULT

The standard asymptotic formulas for, e.g., standard errors, assume that the variables included in the model are *prespecified*

- The "phantom degrees of freedom" induced by having a subjective assessment of the data is not taken into account
- That leads to, e.g., too narrow confidence intervals → too small p-values

Usual practice is to sweep this problem under the rug

One possible solution is too prespecify all variables in advance (including non-linear terms, interactions, ...)

- That is the only way to conserve the advertised coverage rates without further ado

OTHERWISE...

There is plenty of techniques for performing model selections

- Subsetting methods
- Different criteria for comparing models
- Use models that allow for automatic variable selections, e.g.,
 - Random forest
 - Shrinkage

Inference is still a problem, though. Model selection is an active area of research

MODEL SELECTION CRITERIA

SOME MODEL SELECTION CRITERIA

Akaike Information Criterion (1973)

$$\text{AIC}(M) = -2 \log L_M + 2p$$

where L_M is the likelihood of model M and p is the number of parameters

Akaike's rule is to pick the model that minimises the AIC.

Schwarz's criterion (BIC: Bayesian Information Criterion)

$$\text{BIC}(M) = -2 \log L_M + p \log n$$

where n is the sample size

The intuition behind these 2 criteria is that

$$\text{AIC} = \text{bias} + \text{variance}$$

SOME MODEL SELECTION CRITERIA

Mallow's C_p statistic

For a linear model with $p + 1$ parameters

$$C_p = \frac{1}{n} \sum_i (Y_i - \hat{y}_i)^2 + \frac{2}{\hat{\sigma}^2} n(p + 1)$$

Leave-one-out cross-validation

$$\text{LOOCV} = \frac{1}{n} \sum_i \left(Y_i - \hat{y}^{(-i)} \right)$$

where $\hat{y}^{(-i)}$ denote the estimated prediction removing data i .

SUMMARY

There is a theorem which says (roughly)

As $n \rightarrow \infty$, the expected out of sample MSE of the model picked out by LOOCV is close to that of the best model considered.

For n large, LOOCV, AIC and Marlow's C_p will tend to pick the same model

Another theorem says (roughly)

As $n \rightarrow \infty$, if the true model is among those being considered, LOOCV, AIC and C_p will tend to pick models that are strickly larger than the truth

which is ok for prediction

BIC will tend towards choosing the true model. For finite samples, BIC often chooses models that are too simple

MODEL SELECTION STRATEGIES

STEPWISE MODEL SELECTION

Forward stepwise model selection starts with a small model (possibly just containing an intercept)

Consider all possible one-variable expansions of the model

Add the variable that is best according to some criterion

- smallest p-value
- lowest AIC, BIC
- ...

The process is repeated until the criterion stops improving

STEPWISE MODEL SELECTION

Backward stepwise model selection works the other way around

- Start with the largest model you are willing to consider
- Eliminate variables one at a time choose the best variable to remove according to some criterion
- Keep eliminating variables until we no longer improve

Mixed stepwise variable contemplates both adding and removing one variable at each step, and take the best step.

→ it is a greedy algorithm and like trees it is a high variance procedure

SHRINKAGE

The idea of *shrinkage* is to fit a model that contains all p predictors

And use a technique that *constrains* or *regularizes* the coefficients estimates, or equivalently, that *shrinks* the coefficients toward 0

The aim is to reduce the variance

The two best known techniques are *ridge regression* and the *lasso*

RIDGE REGRESSION

Ridge regression shrinks the regression coefficients by imposing a penalty on their size. The ridge coefficients minimise a penalised RSS (in the case of the linear model)

$$\hat{\beta}^{\text{ridge}} = \arg \min_{\beta} \left\{ \sum_{i=1}^n (y_i - \mathbf{x}\beta)^2 + \lambda \sum_{j=1}^q \beta_j^2 \right\}$$

where $\lambda \geq 0$ is a tuning parameter, to be determined separately.

The larger λ , the greater the amount of shrinkage

THE LASSO

The lasso regression also shrinks the regression coefficients towards 0, but unlike ridge regression the penalty in the lasso has the effect of forcing some of the coefficients to be exactly equal to 0

$$\hat{\beta}^{\text{lasso}} = \arg \min_{\beta} \left\{ \sum_{i=1}^n (y_i - \mathbf{x}\beta)^2 + \lambda \sum_{j=1}^q |\beta_j| \right\}$$

Thus the lasso is able to perform variable selection

SHRINKAGE METHODS

One can show that the lasso and ridge regression solve the problems


$$\arg \min_{\beta} \left\{ \sum_{i=1}^n (y_i - \mathbf{x}\beta)^2 \right\} \quad \text{subject to} \quad \sum_{j=1}^q |\beta_j| \leq s$$

and

$$\arg \min_{\beta} \left\{ \sum_{i=1}^n (y_i - \mathbf{x}\beta)^2 \right\} \quad \text{subject to} \quad \sum_{j=1}^q \beta_j^2 \leq s$$

I.e., for every value of λ , there is some s such that the equations above will give the same lasso/ridge coefficient estimates

SHRINKAGE METHODS



SHRINKAGE METHODS

- The tuning parameter λ can be selected via cross-validation

Both methods result in biased regression coefficients estimates, but decreased variance

- The lasso performs variable selection as some coefficients are forced to be equal to 0
- The lasso might perform better if one suspects that many coefficients are close to zero (as might be the case in bioinformatic problems)

The predictors should be standardised, otherwise, possibly important variables with small coefficients (e.g., age) might be dropped.

DIMENSION REDUCTION

The methods discussed so far have controlled variance in two ways

- by using a subset of the original variables
- by shrinking the coefficients towards 0

All these methods have used all the original predictors X_1, X_2, \dots, X_q .

Another possibility is to *transform* the predictors and then fit a model using the transformed variables

- *Principal components regression*: PCA is a technique for reducing the dimension of a $n \times q$ data matrix X
- Partial least squares
- *Propensity score analysis*: = Probability to receive the treatment given the covariates

INFERENCE AFTER MODEL SELECTION

INFERENCE AFTER MODEL SELECTION

It's **BROKEN!!!**

The classical statistical theory assumes that the variables that enter the model are prespecified, i.e., the design matrix is fixed.

The model selection process implies a *random* design matrix. That is not taken into account by the classical theory

Consequences are (usually) too small standard errors → too narrow confidence intervals → too small p-values

INFERENCE AFTER MODEL SELECTION

What works is

- Don't do variable selection
- Make inference on a another data set, e.g,
 1. Divide the data set in two equal parts at random
 2. Build the model in the first part using one of the technique presented before
 3. Make inference in the second part

The obvious drawback is the loss of power for the model building part

- Some complicated (and very recent) resampling based methods
- Model averaging

Automatic techniques for variable prevent you to think!